

Assignment 4

Model Fitting + Stereo/MVS



Line Fitting

- Given a point set with noise and outliers, estimate the parameters: $y = kx + b$
- Implement least-squares solution
- Implement RANSAC (300 iterations)
 1. randomly choose a small subset from the noisy point set ;
 2. compute the least-squares solution for this subset;
 3. compute the number of inliers, if the number exceeds the current best result, update the estimation

Line Fitting

<https://numpy.org/doc/stable/reference/generated/numpy.linalg.lstsq.html>

numpy.linalg.lstsq

`linalg.lstsq(a, b, rcond='warn')`

[\[source\]](#)

Return the least-squares solution to a linear matrix equation.

Computes the vector x that approximatively solves the equation $a @ x = b$. The equation may be under-, well-, or over-determined (i.e., the number of linearly independent rows of a can be less than, equal to, or greater than its number of linearly independent columns). If a is square and of full rank, then x (but for round-off error) is the “exact” solution of the equation. Else, x minimizes the Euclidean 2-norm $\|b - ax\|$. If there are multiple minimizing solutions, the one with the smallest 2-norm $\|x\|$ is returned.


<https://docs.python.org/3/library/random.html>

`random.sample(population, k, *, counts=None)`

Return a k length list of unique elements chosen from the population sequence or set. Used for random sampling without replacement.

Returns a new list containing elements from the population while leaving the original population unchanged. The resulting list is in selection order so that all sub-slices will also be valid random samples. This allows raffle winners (the sample) to be partitioned into grand prize and second place winners (the subslices).

Multi-View Stereo



Network



Source View 2



Reference View



Source View 1

Network









Source View 2



Reference View



Source View 1



Feature

Network






Source View 2




Reference View



Source View 1



Feature



Uniform Depth Samples

Network



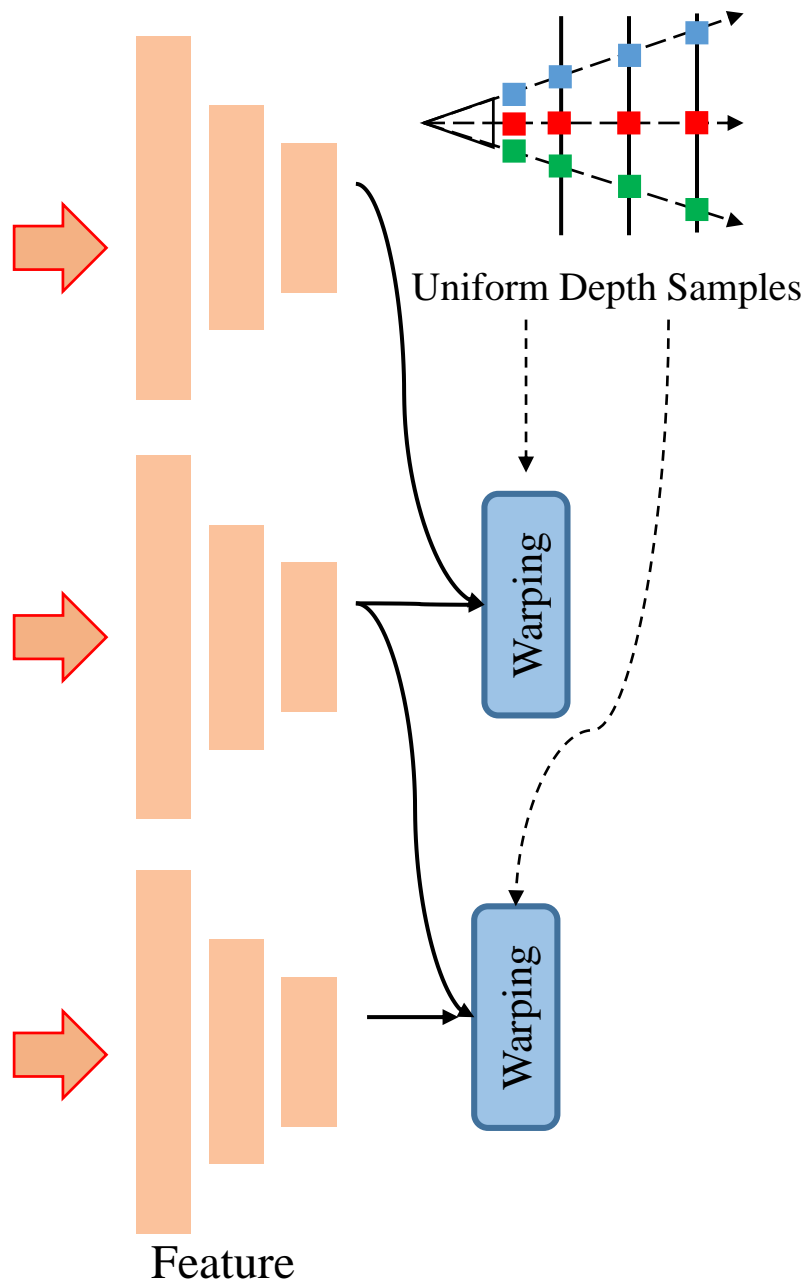
Source View 2



Reference View



Source View 1



Network



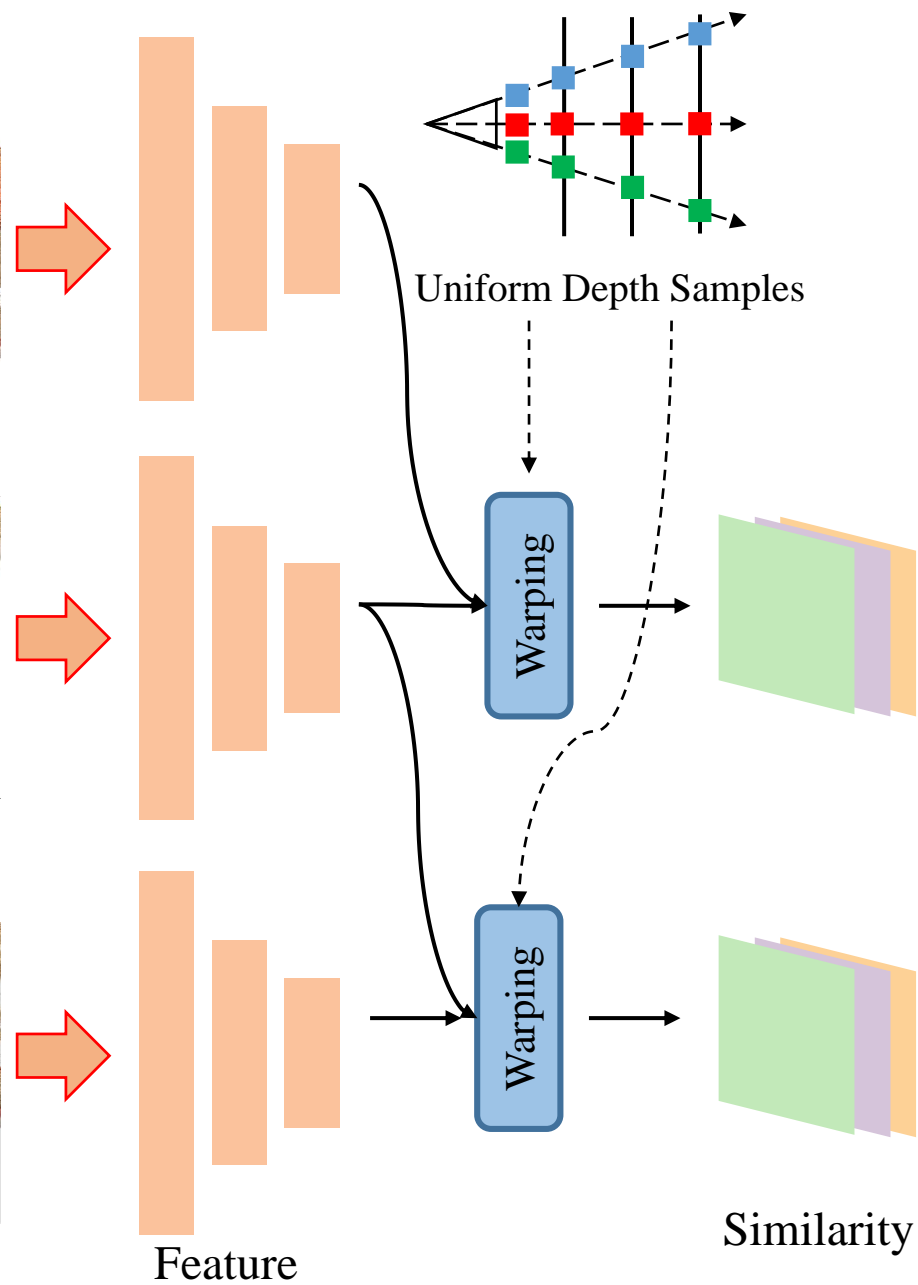
Source View 2



Reference View



Source View 1



Network




Source View 2



Reference View



Source View 1



Network



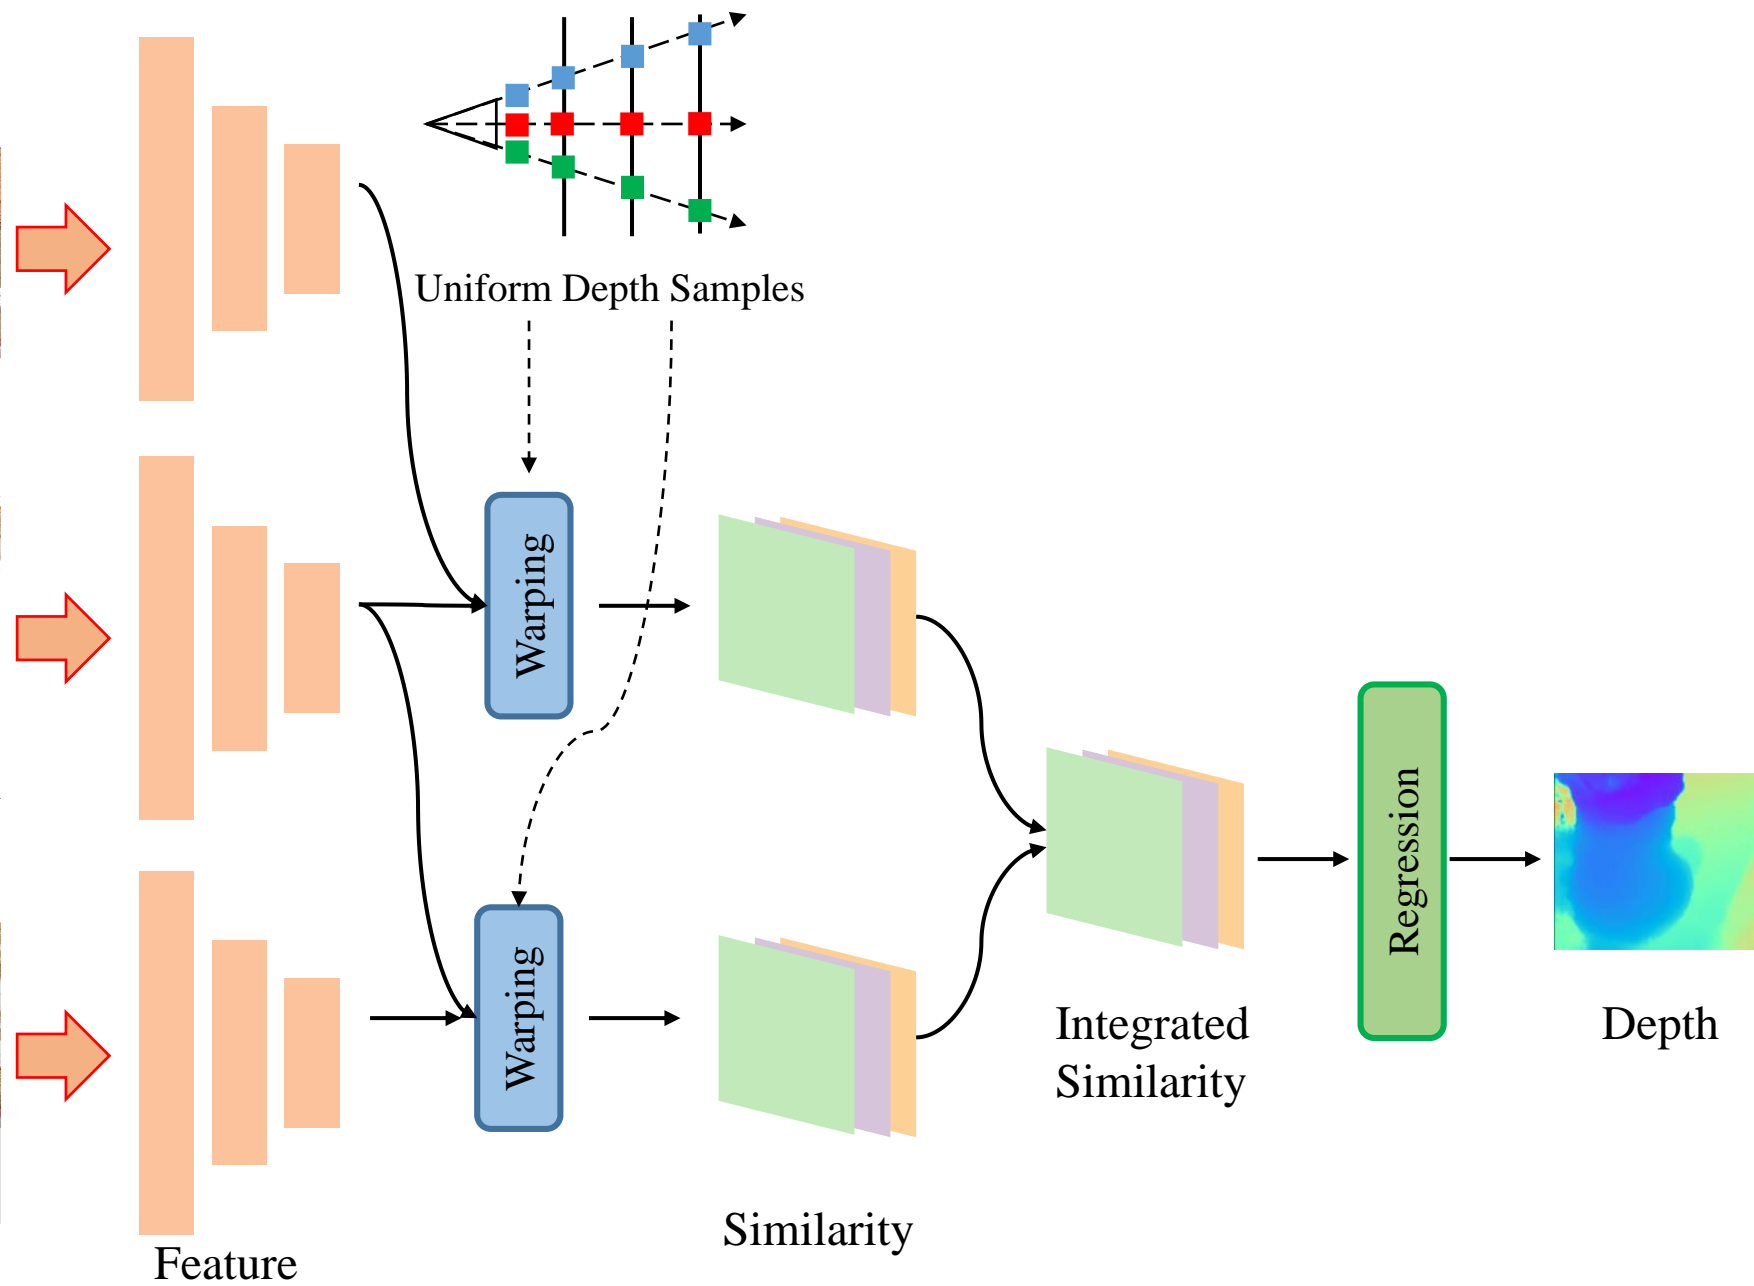
Source View 2



Reference View



Source View 1



Warping

For pixel p in the reference feature and a depth value d_j , we find its corresponding pixel $p_{i,j} := p_i(d_j)$ in source view i , then we get the source feature of $p_{i,j}$

TORCH.NN.FUNCTIONAL.GRID_SAMPLE

```
torch.nn.functional.grid_sample(input, grid, mode='bilinear',  
padding_mode='zeros', align_corners=None) [SOURCE]
```

Given an `input` and a flow-field `grid`, computes the `output` using `input` values and pixel locations from `grid`.

Currently, only spatial (4-D) and volumetric (5-D) `input` are supported.

In the spatial (4-D) case, for `input` with shape $(N, C, H_{\text{in}}, W_{\text{in}})$ and `grid` with shape $(N, H_{\text{out}}, W_{\text{out}}, 2)$, the output will have shape $(N, C, H_{\text{out}}, W_{\text{out}})$.

For each output location `output[n, :, h, w]`, the size-2 vector `grid[n, h, w]` specifies `input` pixel locations `x` and `y`, which are used to interpolate the output value `output[n, :, h, w]`. In the case of 5D inputs, `grid[n, d, h, w]` specifies the `x`, `y`, `z` pixel locations for interpolating `output[n, :, d, h, w]`. `mode` argument specifies `nearest` or `bilinear` interpolation method to sample the input pixels.

https://pytorch.org/docs/stable/generated/torch.nn.functional.grid_sample.html

Dataset

- a part of DTU dataset, a large-scale indoor multi-view stereo dataset
- download link: <https://polybox.ethz.ch/index.php/s/H7SWzSgIQJoSsR2>
- Modify the path in **train.sh** and **eval.sh**



Training and Testing

- Training takes about 9 hours for 4 epochs (cpu-only), If the training process is stopped (e.g. trained 2 epochs), you can resume the training with argument 'resume' (details in train.py)

```
36 parser.add_argument('--loadckpt', default=None, help='load a specific checkpoint')
37 parser.add_argument('--logdir', default='./checkpoints/debug', help='the directory to save checkpoints/logs')
38 parser.add_argument('--resume', action='store true', help='continue to train the model')
39
```

- With the trained model, do depth estimation and point cloud reconstruction (eval.py)

Summary

- Hand in report, images and codes as a zip to moodle before Friday, December 3, 23:59.
- For multi-view stereo, also upload the trained model: model_000003.ckpt

